

拉勾教育

— 互联网人实战大学 —

《Kubernetes 原理剖析与实战应用》

正范

— 拉勾教育出品 —

07 | 有状态应用：Kubernetes 如何通过 StatefulSet 支持有状态应用

Kubernetes 中的另外一种工作负载 **StatefulSet**

主要**用于有状态的服务发布**



在 kubectl 命令行中，一般将 StatefulSet 简写为 sts

在部署一个 StatefulSet 的时候，有个前置依赖对象，即 Service（服务）

```
$ cat nginx-svc.yaml
apiVersion: v1
kind: Service
metadata:
  name: nginx-demo
  namespace: demo
  labels:
    app: nginx
spec:
  clusterIP: None
  ports:
    - port: 80
      name: web
  selector:
    app: nginx
```

```
$ kubectl create ns demo
```

```
$ kubectl create -f nginx-svc.yaml
```

```
service/nginx-demo created
```

```
$ kubectl get svc -n demo
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
nginx-demo	ClusterIP	None	<none>	80/TCP	5s

```
$ cat web-sts.yaml
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: web-demo
  namespace: demo
spec:
  serviceName: "nginx-demo"
  replicas: 2
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
        app: nginx
    spec:
      containers:
        - name: nginx
```

```
replicas: 2
selector:
  matchLabels:
    app: nginx
template:
  metadata:
    labels:
      app: nginx
spec:
  containers:
  - name: nginx
    image: nginx:1.19.2-alpine
    ports:
    - containerPort: 80
    name: web
$ kubectl create -f web-sts.yaml
$ kubectl get sts -n demo
NAME      READY  AGE
web-demo  0/2    9s
```


StatefulSet 的特性

```
$ kubectl get pod -n demo -w
```

NAME	READY	STATUS	RESTARTS	AGE
web-demo-0	0/1	ContainerCreating	0	18s
web-demo-0	1/1	Running	0	20s
web-demo-1	0/1	Pending	0	0s
web-demo-1	0/1	Pending	0	0s
web-demo-1	0/1	ContainerCreating	0	0s
web-demo-1	1/1	Running	0	2s

StatefulSet 的特性

拉勾教育

— 互联网人实战大学 —

```
$ kubectl get pod -n demo -w -l app=nginx
```

StatefulSet 的特性

拉勾教育

— 互联网人实战大学 —

```
$ kubectl get pod -n demo -w -l app=nginx
```

```
$ kubectl get event -n demo -w
```

StatefulSet 的特性

拉勾教育

— 互联网人实战大学 —

```
$ kubectl get event -n demo -w
```

```
$ kubectl scale sts web-demo -n demo --replicas=5  
statefulset.apps/web-demo scaled
```

StatefulSet 的特性

```
$ kubectl get pod -n demo -w
```

NAME	READY	STATUS	RESTARTS	AGE
web-demo-0	1/1	Running	0	20m
web-demo-1	1/1	Running	0	20m
web-demo-2	0/1	Pending	0	0s
web-demo-2	0/1	Pending	0	0s
web-demo-2	0/1	ContainerCreating	0	0s
web-demo-2	1/1	Running	0	2s
web-demo-3	0/1	Pending	0	0s
web-demo-3	0/1	Pending	0	0s
web-demo-3	0/1	ContainerCreating	0	0s
web-demo-3	1/1	Running	0	3s
web-demo-4	0/1	Pending	0	0s
web-demo-4	0/1	Pending	0	0s
web-demo-4	0/1	ContainerCreating	0	0s
web-demo-4	1/1	Running	0	3s

StatefulSet 的特性

```
$ kubectl get event -n demo -w
```

LAST SEEN	TYPE	REASON	OBJECT	MESSAGE
20m	Normal	Scheduled	pod/web-demo-0	Successfully assigned demo/web-demo-0 to kraken
20m	Normal	Pulling	pod/web-demo-0	Pulling image "nginx:1.19.2-alpine"
20m	Normal	Pulled	pod/web-demo-0	Successfully pulled image "nginx:1.19.2-alpine"
20m	Normal	Created	pod/web-demo-0	Created container nginx
20m	Normal	Started	pod/web-demo-0	Started container nginx
20m	Normal	Scheduled	pod/web-demo-1	Successfully assigned demo/web-demo-1 to kraken
20m	Normal	Pulled	pod/web-demo-1	Container image "nginx:1.19.2-alpine" already present on machine
20m	Normal	Created	pod/web-demo-1	Created container nginx
20m	Normal	Started	pod/web-demo-1	Started container nginx
20m	Normal	SuccessfulCreate	statefulset/web-demo	create Pod web-demo-0 in StatefulSet web-demo successful
20m	Normal	SuccessfulCreate	statefulset/web-demo	create Pod web-demo-1 in StatefulSet web-demo successful
0s	Normal	SuccessfulCreate	statefulset/web-demo	create Pod web-demo-2 in StatefulSet web-demo successful
0s	Normal	Scheduled	pod/web-demo-2	Successfully assigned demo/web-demo-2 to kraken
0s	Normal	Pulled	pod/web-demo-2	Container image "nginx:1.19.2-alpine"

StatefulSet 的特性

```
0s      Normal SuccessfulCreate statefulset/web-demo create Pod web-demo-2 in
StatefulSet web-demo successful
0s      Normal Scheduled      pod/web-demo-2      Successfully assigned demo/web-
demo-2 to kraken
0s      Normal Pulled          pod/web-demo-2      Container image "nginx:1.19.2-alpine"
already present on machine
0s      Normal Created          pod/web-demo-2      Created container nginx
0s      Normal Started          pod/web-demo-2      Started container nginx
0s      Normal SuccessfulCreate statefulset/web-demo create Pod web-demo-3 in
StatefulSet web-demo successful
0s      Normal Scheduled      pod/web-demo-3      Successfully assigned demo/web-
demo-3 to kraken
0s      Normal Pulled          pod/web-demo-3      Container image "nginx:1.19.2-alpine"
already present on machine
0s      Normal Created          pod/web-demo-3      Created container nginx
0s      Normal Started          pod/web-demo-3      Started container nginx
0s      Normal SuccessfulCreate statefulset/web-demo create Pod web-demo-4 in
StatefulSet web-demo successful
0s      Normal Scheduled      pod/web-demo-4      Successfully assigned demo/web-
demo-4 to kraken
0s      Normal Pulled          pod/web-demo-4      Container image "nginx:1.19.2-alpine"
already present on machine
0s      Normal Created          pod/web-demo-4      Created container nginx
0s      Normal Started          pod/web-demo-4      Started container nginx
```


StatefulSet 的特性

拉勾教育

— 互联网人实战大学 —

```
$ kubectl scale sts web-demo -n demo --replicas=5  
statefulset.apps/web-demo scaled
```


StatefulSet 的特性

```
web-demo-4 1/1 Terminating 0 11m
web-demo-4 0/1 Terminating 0 11m
web-demo-4 0/1 Terminating 0 11m
web-demo-4 0/1 Terminating 0 11m
web-demo-3 1/1 Terminating 0 12m
web-demo-3 0/1 Terminating 0 12m
web-demo-3 0/1 Terminating 0 12m
web-demo-3 0/1 Terminating 0 12m
web-demo-2 1/1 Terminating 0 12m
web-demo-2 0/1 Terminating 0 12m
web-demo-2 0/1 Terminating 0 12m
web-demo-2 0/1 Terminating 0 12m
```

StatefulSet 的特性

```
0s      Normal SuccessfulDelete statefulset/web-demo  delete Pod web-demo-4 in
StatefulSet web-demo successful
```

```
0s      Normal Killing      pod/web-demo-4      Stopping container nginx
```

```
0s      Normal Killing      pod/web-demo-3      Stopping container nginx
```

```
0s      Normal SuccessfulDelete statefulset/web-demo  delete Pod web-demo-3 in
StatefulSet web-demo successful
```

```
0s      Normal SuccessfulDelete statefulset/web-demo  delete Pod web-demo-2 in
StatefulSet web-demo successful
```

```
0s      Normal Killing      pod/web-demo-2      Stopping container nginx
```

StatefulSet 的特性

```
$ for i in 0 1; do kubectl exec web-demo-$i -n demo -- sh -c 'hostname'; done
```

```
web-demo-0
```

```
web-demo-1
```

StatefulSet 的特性

```
$ kubectl get pod -n demo -l app=nginx -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE
READINESS GATES								
web-demo-0	1/1	Running	0	3h17m	10.244.0.39	kraken	<none>	<none>
web-demo-1	1/1	Running	0	3h17m	10.244.0.40	kraken	<none>	<none>

StatefulSet 的特性

```
$ kubectl run -it --rm --image busybox:1.28 dns-test -n demo
If you don't see a command prompt, try pressing enter.
/ # nslookup web-demo-0.nginx-demo
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: web-demo-0.nginx-demo
Address 1: 10.244.0.39 web-demo-0.nginx-demo.demo.svc.cluster.local
/ # nslookup web-demo-1.nginx-demo
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: web-demo-1.nginx-demo
Address 1: 10.244.0.40 web-demo-1.nginx-demo.demo.svc.cluster.local
```

StatefulSet 的特性

```
$ kubectl delete pod -l app=nginx -n demo
```

```
pod "web-demo-0" deleted
```

```
pod "web-demo-1" deleted
```

```
$ kubectl get pod -l app=nginx -n demo -o wide
```

NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED	NODE
------	-------	--------	----------	-----	----	------	-----------	------

READINESS GATES								
-----------------	--	--	--	--	--	--	--	--

web-demo-0	1/1	Running	0	15s	10.244.0.50	kraken	<none>	<none>
------------	-----	---------	---	-----	-------------	--------	--------	--------

web-demo-1	1/1	Running	0	13s	10.244.0.51	kraken	<none>	<none>
------------	-----	---------	---	-----	-------------	--------	--------	--------

StatefulSet 的特性

```
$ kubectl run -it --rm --image busybox:1.28 dns-test -n demo
If you don't see a command prompt, try pressing enter.
/ # nslookup web-demo-0.nginx-demo
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: web-demo-0.nginx-demo
Address 1: 10.244.0.50 web-demo-0.nginx-demo.demo.svc.cluster.local
/ # nslookup web-demo-1.nginx-demo
Server: 10.96.0.10
Address 1: 10.96.0.10 kube-dns.kube-system.svc.cluster.local

Name: web-demo-1.nginx-demo
Address 1: 10.244.0.51 web-demo-1.nginx-demo.demo.svc.cluster.local
```

如何更新升级 StatefulSet

拉勾教育

— 互联网人实战大学 —

RollingUpdate

是默认的更新策略



OnDelete

当更新策略设置为 OnDelete 时必须手动先删除 Pod 才能触发新的 Pod 更新

StatefulSet 中

支持两种更新升级策略

StatefulSet 的特点:

- 具备固定的网络标记, 比如主机名, 域名等
- 支持持久化存储, 而且最好能够跟实例一一绑定
- 可以按照顺序来部署和扩展
- 可以按照顺序进行终止和删除操作
- 在进行滚动升级的时候, 也会按照一定顺序



Next: 《08 | 配置管理：Kubernetes 管理业务配置方式有哪些? 》

拉勾教育

— 互联网人实战大学 —



关注拉勾「教育公众号」
获取更多课程信息